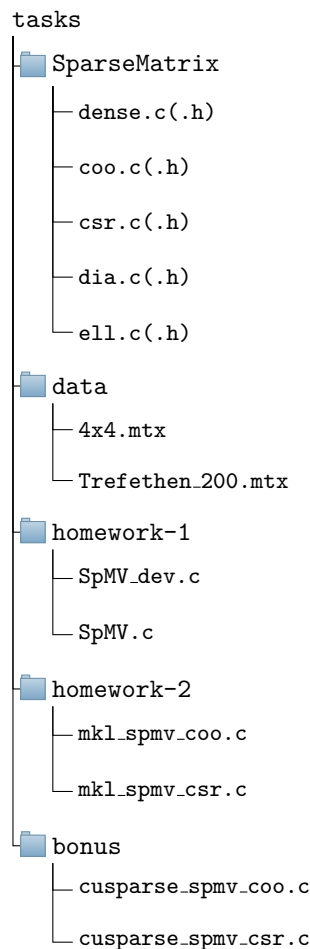


Homework

1. DIRECTORY TREE AND WORKFLOW

The directory tree for the homework is listed as follows:



1.1. **compile the code.** Here is the instruction to compile the code by CMake with the command line tool.

```
1 cd tasks #go into the main directory of homework
2 mkdir build
3 cd build
4 cmake .. #use CMake compile the codes
5 make
```

LISTING 1. How to build the code

1.2. **available matrices.** Several sparse matrices from applications have been provided accompanying with the codes. These matrices can be used for the initial tests. The example listed below is a 4×4 test matrices with 9 entries.

```
1 %%MatrixMarket matrix coordinate real general
2 %-----
3 % TEST
4 %-----
5 4 4 9
6 1 1 1
7 3 1 5
8 1 2 7
9 2 2 2
10 4 2 6
11 2 3 8
12 3 3 3
13 3 4 9
14 4 4 4
```

LISTING 2. Example of MatrixMarket format

The matrices are stored in MatrixMarket format¹. More matrices can be downloaded from the [MatrixMarket](https://math.nist.gov/MatrixMarket/formats.html) or [SuiteSparse Matrix Collection](https://math.nist.gov/MatrixMarket/formats.html).

1.3. **run executables.** Here are examples showing the way to run compiled executables:

```

1 #homework 1
2 ## for development purpose
3 ./homework-1/spmv_dev.exe
4 ## for benchmark purpose
5 ./homework-1/spmv.exe ../data/${NAME_OF_MATRIX_FOR_TEST}
6
7 #homework 2
8 ## MKL with CSR format
9 ./homework-2/mkl_spmv_csr.exe ../data/${NAME_OF_MATRIX_FOR_TEST}
10 ## MKL with COO format
11 ./homework-2/mkl_spmv_coo.exe ../data/${NAME_OF_MATRIX_FOR_TEST}
12
13 #bonus
14 ## cuSPARSE with CSR format
15 ./bonus/cusparse_spmv_csr.exe ../data/${NAME_OF_MATRIX_FOR_TEST}
16 ## cuSPARSE with COO format
17 ./bonus/cusparse_spmv_csr.exe ../data/${NAME_OF_MATRIX_FOR_TEST}

```

LISTING 3. How to run the code

2. HOMEWORK-1

The objective of this homework is to implement SpMV for different sparse matrices format: COO, CSR, DIA and ELL which have been shown in the slides of this module.

For each sparse matrix format, a header file and a related source file have already been implemented within the directory `SparseMatrix`. Below is an example header file `coo.h`, which provides a `struct` for the COO format.

```

1 /*in file SparseMatrix/coo.h*/
2 typedef struct coo
3 {
4     int m; //size of matrix
5     int nnz; //number of non-zero elements in the sparse matrix
6
7     int *rowind; //a pointer storing the row indices of non-zeros elements
8     int *colind; //a pointer storing the column indices of non-zeros elements
9     double *val; //a pointer storing the non-zeros elements
10
11 } coo_t;

```

LISTING 4. COO header file

Similarly, the `structs` for CSR, DIA and ELL can be found within `csr.h`, `dia.h` and `ell.h`.

2.1. **Task 1: sequential SpMV.** This task is to implement sequential SpMV for all available sparse matrix formats.

Requirement:

- CMake version > 3.18
- a C compiler

The implementation should take place with the related source file. for example, for COO format, the sequential SpMV should be implemented in `coo.c` in the `coo_spmv` function.

```

1 /*in file SparseMatrix/coo.c*/
2 void coo_spmv(coo_t *a, double *x, double *y){
3
4     /* put your code here*/
5
6 }

```

LISTING 5. SpMV function in COO source file

The correctness of implementations can be checked by running

```

1 ./homework-1/spmv_dev.exe

```

LISTING 6. How to run the code

It will compare the results obtained for each format with the one of a dense matrix-vector operation. For the validation purpose, the implementation of SpMV should be in the order as: COO, CSR, DIA and ELL. The matrix used for verification is `4x4.mtx` in the directory `data`.

¹<https://math.nist.gov/MatrixMarket/formats.html>

2.2. Task 2: Parallelize SpMV with OpenMP. Try to parallelize the implementations of SpMV for four different sparse matrix storage formats.

Requirement:

- CMake version > 3.18
- a C compiler with OpenMP support

2.3. Task 3: Test with different matrices and threads. The objective of this task is to test the implemented sequential and parallel SpMV with different matrices and different threads number.

Requirement:

- CMake version > 3.18
- a C compiler with OpenMP support

```

1 #Set this environment variable to the number of threads you want to use
2 export OMP_NUM_THREADS=4
3
4 ./homework-1/spmv.exe ../data/${NAME_OF_MATRIX_FOR_TEST}

```

LISTING 7. Run SpMV tests with input matrices and multi-threading

3. HOMEWORK-2

The objective of this homework is to complete the implementation of SpMV based on MKL for both COO and CSR formats. Most parts of implementation have already been available within the directory `homework-2`. You need to fill the missing input arguments when calling the MKL routines.

Requirement:

- CMake version > 3.18
- a C compiler
- Intel MKL (on CLAIIX)

```

1 /*homework-2/mkl_spmv_coo.c*/
2 /*Begin: complete arguments of the following lines*/
3 mkl_sparse_d_create_coo(&mkl_coo_handle, SPARSE_INDEX_BASE_ZERO, , , , , );
4 /*End*/
5
6 /*homework-2/mkl_spmv_csr.c*/
7 /*Begin: complete arguments of the following lines*/
8 mkl_sparse_d_create_csr(&mkl_csr_handle, SPARSE_INDEX_BASE_ZERO, , , , , );
9 /*End*/
10
11 /*Both homework-2/mkl_spmv_coo.c and homework-2/mkl_spmv_csr.c*/
12 /*Begin: complete arguments of the following lines*/
13 mkl_sparse_d_mv(SPARSE_OPERATION_NON_TRANSPOSE, , , , , );
14 /*End*/

```

LISTING 8. Code to be completed in `homework-2/mkl_spmv_csr.c`

The documentation about the three API can be found:

- `mkl_sparse_d_create_coo`
- `mkl_sparse_d_create_csr`
- `mkl_sparse_d_mv`

Attention: For the compilation reason, the build of the source codes of `homework-2` is disabled in default. So if you want to work on this task, please re-run the CMake command as follows with the flag option `-DBUILD_HOMEWORK2=ON`.

```

1 cmake .. -DBUILD_HOMEWORK2=ON
2 make

```

LISTING 9. Re-run CMake to enable the compilation of `homework-2`

4. BONUS: TRY SpMV IN cuSPARSE ON GPU

Try with the SpMV provided by cuSPARSE library on single GPU for both COO and CSR.

Requirement:

- CMake version > 3.18
- a C compiler
- Nvidia CUDA installation (on CLAIIX)

The codes are already completed, the tasks are:

- Run them on supercomputer CLAIIX with different matrices (sparsity pattern, size, etc)
- Identify the time cost of memory transfers between GPU and CPU as a fraction of total time of execution